

Node-screening tests for the ℓ_0 -penalized least-squares problem

Theo Guyard^{*}, Cedric Herzet[†], Clement Elvira[‡]

ICASSP 2022

^{*} Univ Rennes, INSA Rennes, CNRS, IRMAR-UMR 6625, F-35000 Rennes, France

[†] INRIA Rennes-Bretagne Atlantique, Campus de Beaulieu, 35000 Rennes, France

[‡] SCEE/IETR UMR CNRS 6164, CentraleSupélec, 35510 Cesson Sévigné, France

Table of contents

1. The ℓ_0 -penalized least-squares problem
2. Branch-and-bound algorithms
3. Node-screening tests
4. Some numerical results

The ℓ_0 -penalized least-squares problem

Sparse-linear problem

Ingredients of the problem :

- An **observation** $y \in \mathbb{R}^m$
- A **dictionary** $A = [a_i]_{i=1}^n \in \mathbb{R}^{m \times n}$ (columns \equiv **atoms**)

Sparse-linear problem

Ingredients of the problem :

- An **observation** $y \in \mathbb{R}^m$
- A **dictionary** $A = [a_i]_{i=1}^n \in \mathbb{R}^{m \times n}$ (columns \equiv **atoms**)

Objectives :

- Approximate the **observation** as a linear combination of the **atoms**
- The linear combination must be *sparse*

Sparse-linear problem

Ingredients of the problem :

- An **observation** $y \in \mathbb{R}^m$
- A **dictionary** $A = [a_i]_{i=1}^n \in \mathbb{R}^{m \times n}$ (columns \equiv **atoms**)

Objectives :

- Approximate the **observation** as a linear combination of the **atoms**
- The linear combination must be *sparse*

Problem

Find x *sparse* such that $y \simeq Ax$

The vector x weights each atom in the approximation.

ℓ_0 -penalized problem

Idea : Solve the problem

ℓ_0 -penalized least-squares

$$p^* = \begin{cases} \min & \frac{1}{2} \|y - Ax\|_2^2 + \lambda \|x\|_0 \\ \text{s.t.} & \|x\|_\infty \leq M \end{cases} \quad (P)$$

where $\lambda > 0$ is a tuning parameter and M is a big-enough constant.

ℓ_0 -penalized problem

Idea : Solve the problem

ℓ_0 -penalized least-squares

$$p^* = \begin{cases} \min & \frac{1}{2} \|y - Ax\|_2^2 + \lambda \|x\|_0 \\ \text{s.t.} & \|x\|_\infty \leq M \end{cases} \quad (P)$$

where $\lambda > 0$ is a tuning parameter and M is a big-enough constant.

Problem (P) $\xrightarrow{\text{reformulation}}$ Mixed-Integer Program

ℓ_0 -penalized problem

Idea : Solve the problem

ℓ_0 -penalized least-squares

$$p^* = \begin{cases} \min & \frac{1}{2} \|y - Ax\|_2^2 + \lambda \|x\|_0 \\ \text{s.t.} & \|x\|_\infty \leq M \end{cases} \quad (P)$$

where $\lambda > 0$ is a tuning parameter and M is a big-enough constant.

Problem (P) $\xrightarrow{\text{reformulation}}$ Mixed-Integer Program

Properties :

- Quadratic objective
- Linear constraints
- Continuous and integer variables
- Combinatorial problem
- Can be addressed with **Branch-and-Bound (BnB)** algorithms

Branch-and-bound algorithms

Branch-and-bound principle

Idea :

- Enumerate all feasible solutions
 - Use rules to discard irrelevant candidates
- In a nutshell : explore a decision tree and prune uninteresting nodes

Branch-and-bound principle

Idea :

- Enumerate all feasible solutions
- Use rules to discard irrelevant candidates

→ In a nutshell : explore a decision tree and prune uninteresting nodes

Node $\nu = (\mathcal{S}_0, \mathcal{S}_1, \bar{\mathcal{S}})$ where :

- \mathcal{S}_0 : indices of x fixed to zero
- \mathcal{S}_1 : indices of x fixed to non-zero
- $\bar{\mathcal{S}}$: indices not fixed yet

Branch-and-bound principle


Idea :

- Enumerate all feasible solutions
- Use rules to discard irrelevant candidates

→ In a nutshell : **explore a decision tree and prune uninteresting nodes**

Node $\nu = (\mathcal{S}_0, \mathcal{S}_1, \bar{\mathcal{S}})$ where :

- \mathcal{S}_0 : indices of x fixed to **zero**
- \mathcal{S}_1 : indices of x fixed to **non-zero**
- $\bar{\mathcal{S}}$: indices not fixed yet



Question : Does any global solution matches the current constraints ?

Processing node $\nu = (\mathcal{S}_0, \mathcal{S}_1, \bar{\mathcal{S}})$


Question : Does any global solution matches the current constraints ?

Relaxed problem at node ν


$$p_l^\nu = \begin{cases} \min & \frac{1}{2} \|y - Ax\|_2^2 + \frac{\lambda}{M} \|x_{\bar{\mathcal{S}}}\|_1 + \lambda |\mathcal{S}_1| \\ \text{s.t.} & \|x\|_\infty \leq M, x_{\mathcal{S}_0} = 0 \end{cases} \quad (P_l^\nu)$$

Let p_u be an upper bound on p^* . If $p_u < p_l^\nu$, then no optimizers of (P) can match the constraints of node ν .

Exploration and pruning process



Exploration and pruning process



Node-screening tests

Question : Is it possible to detect nodes that cannot yield a global optimizers **without processing them** ?

Dual problem

Question : Is it possible to detect nodes that cannot yield a global optimizers **without processing them** ?

Dual problem at node ν

$$\max_{u \in \mathbb{R}^m} \left\{ D^\nu(u) \triangleq \frac{1}{2} \|y\|_2^2 - \frac{1}{2} \|y - u\|_2^2 - \sum_{i \in \mathcal{S}} [\gamma(a_i^T u)]_+ - \sum_{i \in \mathcal{S}_1} \gamma(a_i^T u) \right\} \quad (D^\nu)$$

Dual problem

Question : Is it possible to detect nodes that cannot yield a global optimizers **without processing them** ?

Dual problem at node ν

$$\max_{u \in \mathbb{R}^m} \left\{ D^\nu(u) \triangleq \frac{1}{2} \|y\|_2^2 - \frac{1}{2} \|y - u\|_2^2 - \sum_{i \in \mathcal{S}} [\gamma(a_i^T u)]_+ - \sum_{i \in \mathcal{S}_1} \gamma(a_i^T u) \right\} \quad (D^\nu)$$

- One common term
- Terms corresponding to the current constraints
- The “pivot” function is defined as $\gamma(t) = M|t| - \lambda$

Direct consequence : The objective of two consecutive nodes differs from one term.

Dual objective link

Direct consequence : The objective of two consecutive nodes differs from one term.

Dual objective link

At node ν , let i be an unfixed index. Then $\forall u \in \mathbb{R}^m$,

$$D^{\nu \cup \{x_i=0\}}(u) = D^\nu(u) + [\gamma(a_i^T u)]_+$$

$$D^{\nu \cup \{x_i \neq 0\}}(u) = D^\nu(u) + [\gamma(a_i^T u)]_-$$

Dual objective link

Direct consequence : The objective of two consecutive nodes differs from one term.

Dual objective link

At node ν , let i be an unfixed index. Then $\forall u \in \mathbb{R}^m$,

$$D^{\nu \cup \{x_i=0\}}(u) = D^\nu(u) + [\gamma(a_i^T u)]_+$$

$$D^{\nu \cup \{x_i \neq 0\}}(u) = D^\nu(u) + [\gamma(a_i^T u)]_-$$

- $\forall u, D^\nu(u) \leq p_i^\nu$: the dual objective can also be used to prune nodes.
- At a given node, we may be able to **prune subnodes without processing them**.

Node-screening test

Node-screening test

Given an upper bound p_u on p^* and a dual point $u \in \mathbb{R}^m$,

$$D^\nu(u) + [\gamma(a_i^\top u)]_+ > p_u \implies \text{Fix } x_i \neq 0 \text{ at node } \nu$$

$$D^\nu(u) + [\gamma(a_i^\top u)]_- > p_u \implies \text{Fix } x_i = 0 \text{ at node } \nu$$

Node-screening test

Node-screening test

Given an upper bound p_u on p^* and a dual point $u \in \mathbb{R}^m$,

$$D^\nu(u) + [\gamma(a_i^\top u)]_+ > p_u \implies \text{Fix } x_i \neq 0 \text{ at node } \nu$$

$$D^\nu(u) + [\gamma(a_i^\top u)]_- > p_u \implies \text{Fix } x_i = 0 \text{ at node } \nu$$

Practical use : If a node-screening test is passed at node ν , one can immediately **fix a new variable** at this node.

Node-screening test

Node-screening test

Given an upper bound p_u on p^* and a dual point $u \in \mathbb{R}^m$,


$$D^\nu(u) + [\gamma(a_i^T u)]_+ > p_u \implies \text{Fix } x_i \neq 0 \text{ at node } \nu$$

$$D^\nu(u) + [\gamma(a_i^T u)]_- > p_u \implies \text{Fix } x_i = 0 \text{ at node } \nu$$

Practical use : If a node-screening test is passed at node ν , one can immediately **fix a new variable** at this node.

Nesting property : If **multiple** node-screening tests are passed, the corresponding variables can be fixed **simultaneously**.

Consequence of passing a node-screening test



Consequence : Less nodes are explored by the BnB algorithm.

Some numerical results

Some numerical results

Synthetic setups :

1. Generate the dictionary randomly (low or high correlation)
2. Generate a k -sparse vector x^*
3. Set $y = Ax^* + \text{noise}$
4. Tune λ and M to (hopefully) recover x^* by solving (P)

Some numerical results

Synthetic setups :

1. Generate the dictionary randomly (low or high correlation)
2. Generate a k -sparse vector x^*
3. Set $y = Ax^* + \text{noise}$
4. Tune λ and M to (hopefully) recover x^* by solving (P)

Methods compared : CPLEX (commercial solver), a tailored BnB and a tailored BnB **with node-screening tests**.

Some numerical results

Synthetic setups :

1. Generate the dictionary randomly (low or high correlation)
2. Generate a k -sparse vector x^*
3. Set $y = Ax^* + \text{noise}$
4. Tune λ and M to (hopefully) recover x^* by solving (P)

Methods compared : CPLEX (commercial solver), a tailored BnB and a tailored BnB **with node-screening tests**.

Corr.	Sparsity	CPLEX		BnB		BnB+scr	
		Nodes	Time	Nodes	Time	Nodes	Time
Low	$k = 3$	16	13.13	19	0.29	15	0.18
	$k = 5$	96	25.89	70	1.5	56	0.75
	$k = 7$	292	60.84	180	5.14	152	3.02
High	$k = 3$	76	1.73	79	0.38	60	0.26
	$k = 5$	1,424	10.18	965	6.39	725	4.18
	$k = 7$	17,647	106.45	10,461	79.29	7,881	52.16